

August 24, 2026

The Honorable Greg Casar
U.S. House of Representatives
Washington, DC 20515

RE: July 30, 2026 Disclosure Regarding Claude Model Behavior in Third-Party Cybersecurity Evaluations

Dear Representative Casar and Members of Congress:

Thank you for your August 10, 2026 letter regarding the cybersecurity evaluation incidents we publicly disclosed on July 30, 2026. Anthropic takes these incidents seriously and believes that the industry as a whole would benefit from a broader conversation about how to evaluate increasingly powerful AI models securely. As discussed below, Anthropic is committed to implementing risk-mitigation measures, and we continue to work with various AI security bodies on shared best practices for the design and security of evaluation environments. That work has been a priority for the company since before these incidents occurred.

Our response provides an overview of our current understanding of the three incidents we disclosed on July 30, based on our review to date and what we are able to share at this stage. Anthropic's review remains ongoing, and we may provide updates as we learn more. Separately, your letter refers to an incident reported by the UK AI Security Institute (UK AISI) on August 4, 2026. That matter is distinct from the three incidents that are the subject of this response, and because our review of it is not yet complete, we address it briefly below.

The Role of Evaluations in Model Development

Cybersecurity evaluations are a critical component of responsible model development. Anthropic evaluates every frontier model before deployment to determine what a model is capable of. These evaluations help determine whether a model can be released to the public and what safeguards need to be developed before it is.

We have run cybersecurity evaluations since February 2025, when we first tested an earlier Claude model against a benchmark of 40 capture-the-flag challenges. Over time, we have increased the number of benchmarks we use, as new ones become available and model

capabilities advance. We frequently run these evaluations with third-party partners who help us develop a wide range of testing scenarios and who can provide an independent assessment of our models' capabilities.

Anthropic's Proactive Identification and Response

On July 21, 2026, OpenAI disclosed that several of its models had escaped an isolated test environment. Beginning on July 23, we reviewed historic evaluation transcripts for evidence of any instances in which a Claude model had reached the internet from within an evaluation. We suspended all internal cybersecurity evaluations that same day, on July 23, after identifying transcripts indicating that Claude may have accessed the internet while undergoing an evaluation conducted by Anthropic's evaluation partner, Irregular. After reviewing 141,006 evaluation runs, Anthropic confirmed three incidents on July 24 in which a Claude model had accessed the internet and gained access to the production infrastructure of three organizations. Together, these incidents involved six evaluation runs in total, four of which impacted the same organization, and two of which impacted two separate organizations.

Anthropic notified its evaluation partner and each of the three affected organizations on July 27, and provided the two affected organizations that responded to our notification with the information necessary for them to assess the impact on their systems. On July 30, Anthropic also notified the security team at PyPI, the public software registry involved in one of the three incidents, and voluntarily notified relevant government authorities in the United States, the United Kingdom, and the European Union. Anthropic published a detailed public account of the three incidents on its blog that same day.

The Three Incidents

The evaluations at issue were running capture-the-flag exercises conducted in an environment built and operated by a third-party evaluation partner. The environment is designed as a private, self-contained network of simulated hosts. The model is given a workstation inside that network and instructed to target and compromise another host on that same network. No host within that network is intended to have a connection to the internet so that nothing outside the simulated exercise can be reached regardless of what the model does within. The model is not provided with any particular method to complete the task, as the purpose of the exercise is to evaluate whether, and how, the model can achieve its goal.

In all three cases, the models were instructed that the environment was a simulation and that they had no access to the internet. However, a misconfiguration in the evaluation environment left the models with a live network path to the internet, allowing the models

to gain unauthorized access to the systems of three organizations. Neither Anthropic nor its evaluation partner was aware of this misconfiguration until the incidents were identified during the July 23 review.

It is important to note that each of these incidents, discussed in further detail below, differed in kind from the Hugging Face incident disclosed by OpenAI. Anthropic's models did not escape their evaluation environment, nor did they circumvent any containment measure intended to keep that environment offline. Rather, because of the misconfiguration, the environment provided internet access when it should not have. Each model had been instructed that its environment was a simulation with no internet access, and each pursued its assigned task on that basis. In each of the six runs, the model did only what it was assigned to do. No model attempted to deliberately escape its environment, disable or evade monitoring, or pursue any objective beyond its assigned task.

Claude Opus 4.7 – The earliest incident occurred in April 2026 during an evaluation run for Claude's Opus 4.7 model. The fictional target company chosen for the exercise happened to share a name with a real, active website domain. Because the model was unable to reach its intended target within the evaluation environment, it discovered that the real organization's systems were reachable instead and, believing this to be the intended target, obtained application and infrastructure credentials and accessed a database containing several hundred rows of production data.

Internal Research Model – A second incident occurred in June 2026 during an evaluation run for an unreleased, internal research test model. In this instance, when the model was unable to reach its intended target, it searched alternatives on the internet. It scanned roughly 9,000 targets on the open internet and ultimately compromised one organization's internet-facing application using basic cyberattack techniques. Notably, later in the run, the model recognized on its own that the system it had reached had no connection to the fictional exercise, concluded the system it had accessed was real, and stopped.

Claude Mythos 5 – A third incident occurred in July 2026 during an evaluation run for Claude's Mythos 5 model. In this instance, the model, operating on the false belief that it was in a simulation, published a software package to PyPI, the public registry from which Python software is distributed. The package was live for approximately one hour and was downloaded and run on 15 real systems, one of which belonged to a security company that routinely scans newly published packages. When that company's scanner ran the package, the model exfiltrated that company's credentials and used them to access the company's infrastructure. The model did not know that company existed when it published the package. Before the model could publish the package, it needed an email address to register an account, and a phone number to create that email address. The model was unable to find a free phone-number service and attempted to obtain funds to pay for one.

We have no indication that any such attempt succeeded. The model ultimately found an email provider that did not require a phone number. At one point, the Mythos 5 model correctly identified that publishing the package would be a real-world attack if its environment were real. It nonetheless convinced itself that it was still in a simulation.

In none of these three incidents were any Anthropic production systems or any Anthropic customer data accessed or affected. In all three incidents, the evaluations were run on infrastructure that is separate from, and has no access to, Anthropic's internal systems or customer data.

Assessment of Claude's Behavior

Based on our findings, these incidents are best understood as a consequence of the misconfiguration, rather than evidence of misaligned goals. In none of these situations did Claude exfiltrate itself or deliberately attempt to escape its test environment. Instead, having reached the internet, the model continued to pursue the specific capture-the-flag task it had been assigned, treating the systems it encountered as part of that task. In each case, the model relied on basic techniques, such as exploiting weak passwords, to access the systems of the affected organizations. None of the incidents involved exploiting any novel, or even advanced, vulnerabilities.

We recognize that a model's stated reasoning is not always a reliable guide to what is driving its behavior, a limitation reflected in our own published research. For that reason, our assessment does not rest on the models' verbalized reasoning in the transcripts alone. We also questioned the models about their beliefs during these episodes, ran resampling experiments that varied elements of the scenario to see how their behavior changed, and conducted further analysis aimed at testing whether a model's stated belief that it was in a simulation matched what was actually driving its behavior.

The models under evaluation retained their standard model-specific safety training. None were "helpful-only" or otherwise altered versions of the underlying model, as is sometimes the case with models intended for research and testing purposes. What they lacked was the additional layer of classifiers and real-time abuse monitoring that we apply to every generally available Claude model. Running an evaluation without those production safeguards is deliberate and standard because the goal is to measure what the underlying model can actually do so that we can build and calibrate the necessary safeguards. That approach is appropriate so long as the evaluation is contained. To be clear, the production safeguards would have prevented the behavior identified in these three incidents. The configuration that resulted in these incidents (i.e., no production safeguards in place, an explicit instruction to attack, and an environment with unintended internet access) does not exist in any generally released product.

Anthropic's Ongoing Response

Anthropic temporarily paused all internal cybersecurity evaluations on July 23. Since then, it has resumed cybersecurity evaluations of pre-release models, subject to the enhancements described below. Anthropic is also in the process of engaging METR, an independent AI evaluation organization, to conduct a third-party review of the incidents. Although METR will have access to transcripts as part of its review, we are not in a position to publicly release the full transcripts given our ongoing review and in order to protect the affected organizations.

Anthropic is implementing a number of enhancements to its evaluation program, including:

- Improving processes for validating the network configuration of our evaluation environments;
- Deploying a monitor that reviews each cyber evaluation as it runs in real time, and blocks and escalates any action that falls outside the scope of the evaluation;
- Migrating our cyber evaluation infrastructure to a more hardened isolation platform;
- Requiring that either Anthropic or the evaluation partner takes clear responsibility for monitoring the run, with each partner required to confirm in writing that a defined set of isolation, validation, and monitoring practices is in place before testing begins or continues; and
- Revising how we write our evaluation prompts so that a model is told explicitly which systems are within scope (a change that, on its own, would likely have prevented these incidents).

In addition to the above, Anthropic is also working with government AI security bodies on shared standards for the design and security of evaluation environments, which is consistent with the position we have taken publicly on the need for a robust third-party testing ecosystem.

Our review of the evaluations in which these three incidents arose remains ongoing.

UK AISI Incident

Your letter also refers to the incident disclosed by the UK AISI on August 4 relating to testing conducted on two models, one of which was a Claude model. The UK AISI conducted that evaluation independently of Anthropic and has published its own findings. According to UK AISI's report, the models engaged in potentially harmful activity as a result of what it described as "deliberately permissive conditions" (i.e., conditions in which safeguards had been removed, internet access provided, and no restrictions on how the

internet should be used). UK AISI found no evidence that the models escaped from a secure environment.

We have not yet received the underlying materials from the evaluation to conduct our independent review of the incident. Accordingly, we are not able to characterize the findings beyond what UK AISI has itself published. We are actively working with UK AISI to obtain those materials, and would be glad to share our assessment with your staff once our review is complete.

Closing

We identified the three incidents disclosed on July 30 independently, notified the affected organizations, and provided a detailed public account of what happened. We hope the information set out above is helpful in your review. Should your staff find it useful, we would welcome the opportunity to discuss these incidents further, and to take part in the broader conversation we think the industry should have about how to securely evaluate increasingly powerful AI models. Anthropic values its engagement with Congress and remains committed to working cooperatively with you and your colleagues.

Sincerely,

Anthony Cimino
Head of US Federal Affairs
Anthropic PBC